


Restaurant Recommender

Daniil Abbruzzese



Agenda



01

Introduction

My motivation for selecting my idea

02

Process

How I created my data product

03

Live Demo

Generating restaurant recommendations

04

Next steps

Incorporating content based filtering



The Problem

Choosing which restaurant to eat at can be **overwhelming** due to the sheer amount of choices.

Googling “best restaurants in my city,” yields too many **generic results** that are not personalized to the individual

As a result, I end up going to the same few restaurants that I discovered through word of mouth.

Solution: personalized recommendations

I have created a collaborative filtering model, trained on the **Yelp Open Dataset**, that delivers **personalized restaurant recommendations** to users.

My recommendation system learns what users like and dislike in a restaurant, whether or not they know it themselves. It thus helps users **expand their horizon**.



Process




Preparing that Data

Filtering my data and
generating a user-item
matrix



Model Building


Selecting a collaborative
filtering model and tuning
hyper parameters



Recommendations

Despite being red, Mars is
actually a cold place. It's
full of iron oxide dust





Preparing the Data





Preparing the Data

1. Filters: "Restaurants" in California only
2. Merging reviews dataset with business dataset
3. Users who left at least 5 reviews
4. Averaging ratings for users who left multiple reviews at same restaurant
5. Pivoting the df to get a UI matrix



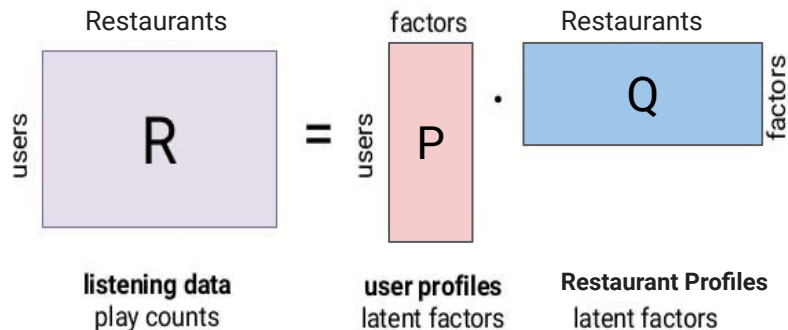
Model Building



Matrix Factorization with Stochastic Gradient Descent

Matrix factorization

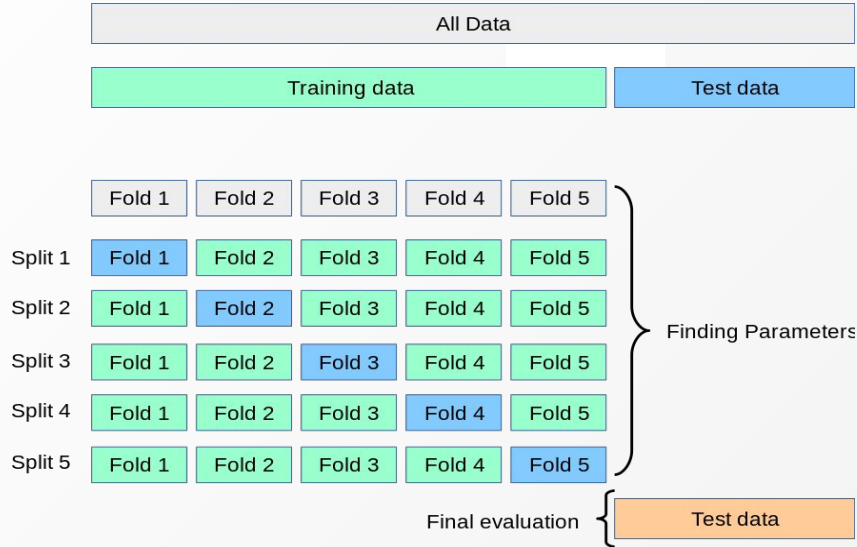
Model listening data as a product of latent factors




15

- Goal is to discover relationship to “Latent Factors” for users and restaurants
- Represented by Matrices P and Q
- Taking the dot product of P and Q yields an estimate of R which minimizes the loss
- Train / test split: for each user, I removed at least 1 rating for testing
- Compares predicted rating to actual to calculate the loss, adjusts values in the direction opposite to the loss function

Hyper Parameter Tuning




- Used cross validation to identify ideal hyper parameters
- Hyper parameters:
 - K: # Latent features
 - Learning Rate: size of the steps taken during gradient descent optimization
 - Regularization term: penalizes the magnitude of the parameters to avoid overfitting
 - Number of iterations: Number of times model updates values



Bayesian Optimization

- I found that grid search was taking too much time
- Used Bayesian Optimization to guide the search for best hyper parameters
- Uses a probabilistic model to predict which set of hyper parameters will perform best on the validation set based on past results
- Reduces number of evaluations needed by focusing on areas in grid space with highest potential for improvement, making it much faster



Recommendations



Delivering Recommendations



1. Retraining model on all data
2. User inputs favorite restaurants in a list
3. List is converted to a user ratings vector (row in UI matrix)
4. A user latent feature vector is derived given our restaurant latent feature matrix (Q)
5. We predict ratings across all restaurants by taking dot product of user's latent feature vector and restaurant latent feature matrix
6. Sort by 10 highest ratings, giving recommendation

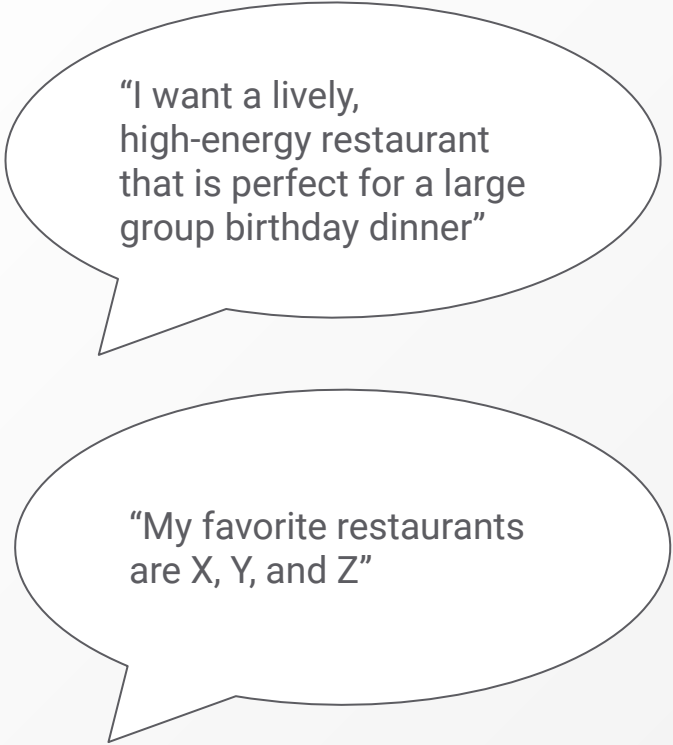


Live Demo



Next Steps: Hybrid Model





"I want a lively,
high-energy restaurant
that is perfect for a large
group birthday dinner"

"My favorite restaurants
are X, Y, and Z"

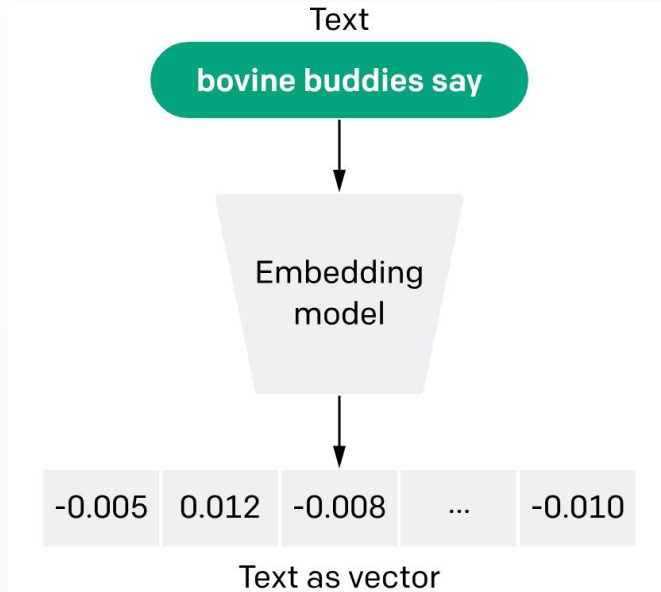
Hybrid Model

Incorporate **content-based** filtering into the recommendation system.

Users can filter for the **type of atmosphere** they want (Date night, birthday dinner, calm vs lively atmosphere, etc).

Open AI's embeddings can help with this. Embeddings turns text into a high dimensional vector representation that captures semantic meaning

High-level plan



1. New user types in their desired atmosphere
2. Convert this input into an embedding
3. Compute the cosine similarity between the user input embedding and the embeddings of various restaurants
4. Rank Restaurants: Rank the restaurants based on their similarity score with the user's input.
5. Hybrid model integration: of the top 30 restaurants that match the desired atmosphere, recommend top 5 for predicted score from collaborative filtering model

Obtaining one embedding per restaurant

Approach one:

- Filter for reviews that have keywords relating to ambiance
- Generate embeddings for each, take the average
- **Downside:** will miss out on reviews that don't contain keywords but still convey information on ambiance

Approach two:

- Use GPT 3 or 4 to read all reviews and generate a single summary description that captures ambiance
- Convert summary caption into an embedding
- **Downside:** more timely and costly



Thank You
